

Class Size: Answers through Wisdom and Justice

The Question

The Tennessee STAR experiment (Mosteller 1995) randomly assigned more than 6,000 students in about 80 Tennessee schools to small classes (13–17 students), regular classes (22–25), or regular classes with a teacher’s aide, from kindergarten through third grade. Our data records each student’s kindergarten class-type assignment, race, years spent in small classes across grades K through 3, and fourth-grade score on the math computation portion of the Stanford Achievement Test.

Imagine you work for the Texas Department of Education. You want to understand the effect of small classes on student performance in Dallas. In Texas, there is data available like the data in the STAR project.

The Texas Department of Education asks: What is the average causal effect of assigning a Dallas student to a small class on their fourth-grade math score?

Wisdom

Preceptor Table

What are the units, outcome, and covariates for this question?

- Units:

Each row is one Dallas student currently enrolled in grades K through 3, whose fourth-grade math score will arrive over the next four years. The Department cares about the causal effect for all of them, so the Preceptor Table includes rows for next year's fourth graders, the fourth graders the year after that, and so on — four cohorts, tens of thousands of students.

- Outcomes:

Three *potential* outcomes per row — the math score if assigned to a small class, if assigned to a regular class, and if assigned to a regular class with an aide. The scenario names the treatment — small classes — but not the control: the comparison might be against regular classes, regular classes with an aide, or both, so the Preceptor Table must carry all three potential outcomes. The Preceptor's Question leaves the control unspecified too — on purpose. “Assigned to a small class” compared to what? Small versus regular and small versus regular-with-aide are two *different* causal effects, and both are worth considering.

- Covariates:

None are forced by the question, though race and years in small classes will earn their way into the model.

- Treatment (if any):

The *kindergarten* class-type assignment — small versus regular versus regular with aide — which STAR randomized at school entry. With three levels, each student has three potential outcomes, and each pairwise contrast is a causal effect.

- Quantity of interest:

The average difference between the potential outcome under a small-class assignment and the potential outcome under a control assignment. The question does not say which control — regular, or regular with an aide — so there are two causal effects on the table, and the Department should look at both.

Describe the key components of a Preceptor Table for this problem; Use words like “units,” “outcomes,” and “covariates.”

The Preceptor Table should have one row per Dallas student now in grades K through 3 — four cohorts of future fourth graders. The columns are Student, the three potential outcomes (Score if Small, Score if Regular, Score if Regular with Aide), and the Treatment (Class Assignment).

Preceptor Table --- Texas Department of Education (causal)¹

Unit ²	Potential Outcomes ³			Treatment ⁴
Student	Score if Small	Score if Regular	Score if Regular with Aide	Class Assignment
Ethan Williams	712	705	707	Small
Ava Martinez	703	698	700	Regular
...	...	...	...	...
Noah Thompson	721	716	718	Regular with aide

¹ If all the information in this table were available, we could answer the question: What is the average causal effect of assigning a Dallas student to a small class on their fourth-grade math score?

² Each row represents one Dallas student currently in grades K through 3 — four cohorts of future fourth graders, tens of thousands of students.

³ The three potential outcomes: the student's fourth-grade math computation score under each possible kindergarten class assignment. With a three-level treatment, each student has two unobservable potential outcomes — the cross-hatched cells — because the student actually experienced only one assignment.

⁴ The kindergarten class-type assignment: Small (13-17 students), Regular (22-25), or Regular with aide. This is the variable whose causal effect the question is about.


Exploratory Data Analysis

Create a plot that shows the distribution of `g4math`.

```
star |>
  drop_na(g4math) |>
  ggplot(aes(x = g4math)) +
  geom_histogram(bins = 30, fill = "steelblue", color = "white") +
  labs(
    title = "Fourth-Grade Math Scores",
    subtitle = "Scores cluster around 700 -- and most students have no score at all",
    x = "Fourth-grade math score",
    y = "Number of students"
  ) +
  theme_minimal()
```

Fourth-Grade Math Scores


Scores cluster around 700 -- and most students have no score at all



Scores are roughly bell-shaped around 709 — but the subtitle carries the real news: fewer than 40% of the students in the experiment have a fourth-grade math score recorded. Before trusting any model, we need to know whether that missingness is related to anything we care about; reassuringly, the share missing is nearly identical — about 62% — in all three class types, though balance cannot rule out bias if the *kinds* of students who go missing differ.

Create a plot that shows the relationship between `g4math` and class type.

```
star |>
  drop_na(g4math) |>
  ggplot(aes(x = kind, y = g4math, fill = kind)) +
  geom_boxplot(width = 0.5, alpha = 0.7, show.legend = FALSE) +
  labs(
    title = "Math Scores by Class Type: Boxplot",
    subtitle = "The three class types look almost identical",
    x = "Class type",
    y = "Fourth-grade math score"
  ) +
  theme_minimal()
```



The three class types are visually indistinguishable. If small classes help fourth-grade math, the effect is too small to see with the naked eye — a useful calibration before the modeling in Courage.

Justice

Population Table

Describe the key components of a Population Table for this problem.

One row per student per school year, drawn from the population of US elementary students from roughly 1985 to 2030, with the causal structure explicit. The Data rows are Tennessee STAR students from 1985; the Preceptor rows are Dallas students now in grades K through 3, whose fourth-grade scores arrive from 2027 through 2030. The columns are Source, Student, Year, the three potential outcomes, and the Treatment.

Population Table --- Texas Department of Education (causal) ¹						
Source	Unit/Time ²		Potential Outcomes ³			Treatment ⁴
	Student	Year	Score if Small	Score if Regular	Score if Regular with Aide	Class Assignment
...	...	...	...	...	...	...
Data	Dorothy Baker	1985	704	...	...	Small
Data	James Whitfield	1985	...	691	...	Regular
Data	...	...	...	...	...	...
Data	Angela Foster	1985	...	...	715	Regular with aide
...	...	...	...	...	...	...
Preceptor	Ethan Williams	2027	712	705	707	Small
Preceptor	Ava Martinez	2028	703	698	700	Regular
Preceptor	...	...	...	...	...	...
Preceptor	Noah Thompson	2030	721	716	718	Regular with aide
...	...	...	...	...	...	...

¹ This table combines the Tennessee STAR data with the Texas Department of Education's Dallas students today, both drawn from the population of US elementary students from roughly 1985 to 2030.

² Data rows are STAR students in Tennessee, 1985. Preceptor rows are Dallas students now in grades K through 3, whose fourth-grade scores arrive from 2027 through 2030. The gaps of four decades and two states are the stability and representativeness concerns.

³ In Data rows, only the potential outcome matching the student's actual assignment was observed; the other two are shown as '...' because no one measured them. In Preceptor rows, all three potential outcomes are filled in with the truth, and the two unobservable ones are cross-hatched.

⁴ In Data rows, the kindergarten class assignment was randomized by the experimenters at school entry — the basis for unconfoundedness. In Preceptor rows, it is whatever assignment Dallas schools would give, which nothing guarantees to be random.

Assumptions

Counter-example to the assumption of validity:

The data's `g4math` column is the math *computation* portion of a 1980s Stanford Achievement Test; a Dallas student today takes a different test, with different content and scaling — and their district's "math score" is likely a total score, not a computation subscore. Likewise "small class" meant 13–17 students in STAR — if Dallas's version of "small" is 20, the treatment column in the data and the treatment column in the Preceptor Table share a name but not a meaning.

Counter-example to the assumption of stability:

Students' lives have changed since 1985 — screens, curricula, teacher training, post-pandemic learning loss. If small classes helped in 1985 mainly by increasing teacher attention per student, and today's classrooms distribute attention differently (aides, technology, tutoring software), the coefficient estimated on 1985 data is not the coefficient for 2026-27.

Counter-example to the assumption of representativeness:

Tennessee schools that volunteered for STAR in 1985 are not a random draw from anything — and neither are today's Dallas students, whose demographics, funding, and family backgrounds differ sharply from 1980s Tennessee. If small classes help disadvantaged students more (as the original study suggested), differences in student mix between Tennessee then and Dallas now change the answer.

Counter-example to the assumption of unconfoundedness:

Class type was randomly assigned, so unconfoundedness should hold by design. The honest worries are execution and erosion: administrators may have quietly moved favored children into small classes, and over four years students did move between class types — so the *initial* kindergarten assignment — the treatment our data records — is random, but the class a student actually sat in by third grade is not guaranteed to be.

Probability Family and Link Function

What probability family should we use?

Normal Distribution

Why should we use that probability family?

The Normal distribution is the standard choice for continuous outcomes, and a math score is continuous.

What is the link function for our outcome?

The identity link. Our model will be a linear regression.
